

Does the Model Know Why?

A Failure-Mode Cross-Checking Protocol for Verifying SHAP Explanations in Multi-Mechanism Predictive Maintenance

Koshik Debanath
Independent Researcher
koshik.debanath@gmail.com

Abstract

Post-hoc explanation methods such as SHAP are routinely applied to predictive-maintenance classifiers, but the resulting attributions are rarely checked against anything beyond the analyst’s intuition. We use a property of the AI4I 2020 predictive maintenance dataset that is usually discarded: its binary **Machine failure** label is the logical OR of five independently defined failure mechanisms (tool wear, heat dissipation, power, overstrain, and a random failure with no physical cause), each with an exact, published triggering rule. We train an imbalance-aware gradient-boosted classifier on the composite label only – withholding the five sub-labels from the feature set entirely – and then use the sub-labels purely as a post-hoc audit: for each mechanism, we compute the model’s mean absolute SHAP attribution restricted to the rows where that mechanism fired, and check whether the top attributed feature matches the mechanism’s known physical cause. The random-failure mechanism, which by construction has no learnable predictor, serves as a negative control. Across five mechanisms, our tuned XGBoost model (precision 1.00, recall 0.82, F1 0.90 on the held-out failure class) attributes each real mechanism to its correct physical cause and shows no dominant feature for the random mechanism. We argue this cross-checking protocol – using a dataset’s own latent causal substructure as a faithfulness probe, rather than constructing synthetic ground truth for the purpose – generalizes to any predictive-maintenance setting where a coarse target is known to aggregate several distinct, partially-labeled failure modes, which is the common case in industrial practice.

1 Introduction

Predictive maintenance (PdM) models are increasingly built with gradient-boosted trees or neural networks and explained after the fact with SHAP [Lundberg and Lee, 2017]. The explanation is typically presented as evidence that the model “makes sense” – a beeswarm plot, a bar chart of global feature importance – but is rarely checked against an independent source of truth about *why* the underlying process actually failed. This is a meaningful gap: a model can achieve strong discriminative performance while keying off a spurious correlate of failure rather than the mechanism itself, and a plausible-looking SHAP plot will not by itself reveal the difference [Vollert and Theissler, 2024].

The AI4I 2020 dataset [Matzka, 2020] offers an unusual opportunity to close this gap without constructing anything artificial for the purpose. Its target, **Machine failure**, is by design the logical OR of five independent sub-conditions – tool wear failure (TWF), heat dissipation failure (HDF), power failure (PWF), overstrain failure (OSF), and a random failure (RNF) that fires with fixed probability regardless of process state. Each of the first four has an exact, published closed-form trigger (Table 1); RNF has none. Critically, a model trained to predict **Machine failure**

is never given these five labels – it only ever sees the OR of them. This is exactly the situation faced in real industrial settings, where a single “needs maintenance” flag is often an aggregate of several distinct root causes that are logged separately (or not logged at all) but never exposed to the model as a training signal.

This gives us a natural faithfulness probe: after training on the composite label alone, we can go back and ask, for the subset of rows where a specific mechanism actually fired, whether the model’s SHAP attributions concentrate on the feature that mechanism’s own definition implicates. If they do not, that is a warning sign that the classifier is exploiting a different, possibly spurious, correlate. If RNF – which has no real cause – shows no dominant attributed feature, that is a sanity check that the method is not simply finding *something* to blame in every subgroup by construction.

Contributions.

1. An imbalance-aware benchmark of seven model families (Logistic Regression, SVM, k -NN, Decision Tree, Random Forest, XGBoost, and a small PyTorch MLP) on AI4I 2020, with a resampling-technique bake-off (SMOTE vs. Borderline-SMOTE vs. ADASYN), hyperparameter and decision-threshold tuning, and 5-fold cross-validation.
2. A lightweight cross-checking protocol that uses a dataset’s own latent, partially-labeled causal substructure – rather than purpose-built synthetic ground truth – to audit whether a trained model’s SHAP explanations are faithful to the real generating mechanism, including a built-in negative control.
3. An empirical demonstration that a well-tuned XGBoost classifier passes this audit on all five mechanisms of AI4I 2020, and a discussion of what a failed audit would look like and how the protocol generalizes to other multi-mechanism PDM datasets.

Table 1: The five failure mechanisms underlying **Machine failure** in AI4I 2020 [Matzka, 2020]. Only the composite OR label is ever shown to the model; the individual triggers are withheld from training and used only for the audit in Section 4.5.

Code	Name	Trigger condition	Count
TWF	Tool wear failure	tool wear falls in a randomly assigned 200–240 min window	46
HDF	Heat dissipation failure	(process temp. – air temp.) < 8.6 K and rpm < 1380	115
PWF	Power failure	torque $\times$ rpm (rad/s) outside [3500 W, 9000 W]	95
OSF	Overstrain failure	tool wear $\times$ torque exceeds a type-specific limit (11/12/13k)	98
RNF	Random failure	fixed 0.1% chance, independent of all process variables	19

2 Related Work

Matzka [2020] introduced the AI4I 2020 dataset alongside a LIME-based illustration of local explanations for individual predictions, but did not systematically audit attributions against the five

sub-mechanisms across the dataset. Several subsequent works apply SHAP to AI4I 2020 for global feature-importance reporting – e.g., a recent milling-machine study [Zhao et al., 2025] reports SHAP summary plots for tree ensembles trained on this dataset – but, to our knowledge, none of them cross-check per-mechanism attributions against the known triggering rules, nor use the random-failure mechanism as a negative control. A recent physics-informed benchmarking study [Tseng, 2026] constructs engineered features from the same sensor variables and reports a Random Forest positive-class F1 of 0.722 on AI4I 2020, and a connected-vehicles study [Khemani and Qureshi, 2026] benchmarks LightGBM under a stricter SMOTE-confined-to-training-folds cross-validation protocol, reporting macro-F1 0.814 ± 0.012 . We compare directly against both in Section 5.

Dhande et al. [2026] take a related but distinct angle: they build formal causal structural models on a similarly-sized CNC failure dataset to optimize maintenance decisions under asymmetric cost, and show causal features generalize better than raw correlational ones. Our goal is narrower and complementary – we do not build a causal model of the process; we audit whether an ordinary correlational classifier’s *post-hoc explanation* happens to align with the true mechanism, using labels that already encode that mechanism.

Work on faithfulness evaluation for explainability methods generally relies on *synthetic* ground truth: data is generated from a known functional form specifically so that the correct attribution is computable in closed form [Liu et al., 2021]. Our approach differs in where the ground truth comes from – rather than designing a synthetic task for the purpose of evaluating explanations, we reuse sub-labels that already exist in an otherwise ordinary, publicly available PdM benchmark, aggregated away in the task the model is actually asked to solve. AI4I 2020 is itself a simulated dataset generated from these same closed-form rules, a caveat we return to in Section 6.2.

3 Dataset

AI4I 2020 [Matzka, 2020] contains 10,000 rows and 14 columns: a row identifier and product ID (dropped), a categorical product quality variant (**Type** $\in \{L, M, H\}$, 50/30/20% of rows respectively), five continuous/integer process variables (air temperature, process temperature, rotational speed, torque, tool wear), the composite **Machine failure** label, and the five mechanism indicators of Table 1. **Machine failure** is 1 for 339 of 10,000 rows (3.39%), a severe class imbalance typical of real PdM data.

4 Methodology

4.1 Feature engineering

We construct four features directly from the mechanism definitions in Table 1, without using the mechanism labels themselves: *Power* = torque $\times$ rpm (rad/s) (mirrors the PWF rule), *Temp-diff* = process temperature – air temperature (mirrors HDF), *Overstrain* = tool wear $\times$ torque (mirrors OSF), and *Overstrain-ratio* = Overstrain normalized by the type-specific limit. This gives the model, in principle, direct access to each mechanism’s driving quantity without ever seeing which mechanism fired.

4.2 Leakage-safe feature extraction

UID and **Product ID** are identifiers and are dropped. TWF, HDF, PWF, OSF, and RNF are components of the target by construction (**Machine failure** = $\bigvee_i \text{mode}_i$) and are excluded entirely from the feature matrix; they are retained only as held-out audit labels.

4.3 Split, resampling, and models

We use a stratified 70/20/10 train/validation/test split. The preprocessor (standard-scaling + one-hot encoding) is fit on the training split only. We compare SMOTE [Chawla et al., 2002], Borderline-SMOTE [Han et al., 2005], and ADASYN [He et al., 2008] by resampling the training split and scoring a fixed default XGBoost [Chen and Guestrin, 2016] on the untouched validation split; the winner (Borderline-SMOTE in our run) is used for the remainder of the pipeline. We train seven model families – Logistic Regression, SVM (RBF), k -NN, Decision Tree, Random Forest, XGBoost, and a small PyTorch feed-forward network – on the resampled training data and evaluate on the untouched, imbalanced validation and test splits.

4.4 Tuning and validation

The best baseline model (XGBoost) is tuned with a 40-iteration, 5-fold stratified `RandomizedSearchCV` scored on F1. Its decision threshold is then chosen from the precision-recall curve on the validation split, maximizing F1, rather than using the default 0.5 cutoff. We additionally run 5-fold stratified cross-validation, refitting preprocessing, resampling, and the tuned hyperparameters inside each fold, to obtain a variance estimate rather than relying on a single 1,000-row test split. We check probability calibration with isotonic regression and the Brier score before trusting the tuned model’s raw scores for threshold selection.

4.5 XAI cross-checking protocol

This is the core methodological contribution. We compute SHAP `TreeExplainer` values for the final tuned model over the full dataset (not just the small test split – several mechanisms have fewer than 20 total positive rows, too few to analyze reliably within a 1,000-row test partition alone; this uses no additional training signal, only post-hoc explanation of an already-fixed model). For each mechanism m and each feature j , we compute

$$I_{m,j} = \frac{1}{|S_m|} \sum_{i \in S_m} |\phi_j(x_i)|,$$

where S_m is the set of rows for which mechanism m fired and $\phi_j(x_i)$ is the SHAP value of feature j for row i . For each real mechanism (TWF, HDF, PWF, OSF) we check whether $\arg \max_j I_{m,j}$ matches the feature(s) implicated by that mechanism’s known trigger. For RNF, which has no real cause, we instead check that no single feature dominates: we compute the share of total importance held by the top feature, $\max_j I_{\text{RNF},j} / \sum_j I_{\text{RNF},j}$, and treat a share below 0.35 as passing – i.e., the audit is designed so that a model that spuriously “explains” random failure would fail it.

5 Results

5.1 Model comparison

Table 2 reports positive-class (`Machine failure`=1) metrics on the held-out test split. XGBoost dominates the comparison even before tuning; tree ensembles substantially outperform the linear and instance-based models, consistent with the class boundary being nonlinear and threshold-based (Section 5.3) rather than a simple linear separation. Hyperparameter and threshold tuning lift positive-class F1 from 0.87 to 0.90, driven almost entirely by a precision gain ($0.86 \rightarrow 1.00$) at a moderate recall cost ($0.88 \rightarrow 0.82$).







Figure 1: Test-set performance profile of the winning model. (a) Positive-class F1 across all classical models (tuned XGBoost highlighted). (b) Confusion matrix for the tuned model at its selected threshold. (c) ROC and (d) precision-recall curves for all classical models, tuned XGBoost highlighted. The PyTorch MLP is omitted from (a), (c), (d) for legibility – its numbers are in Table 2.

Table 2: Positive-class (`Machine failure=1`) test-set metrics, sorted by F1.

Model	Precision	Recall	F1	ROC-AUC
XGBoost (tuned)	1.00	0.82	0.90	0.984
XGBoost (baseline)	0.86	0.88	0.87	0.988
Random Forest	0.79	0.88	0.83	0.989
Decision Tree	0.68	0.82	0.75	0.927
PyTorch MLP	0.53	0.91	0.67	0.978
SVM (RBF)	0.42	0.88	0.57	0.976
k -NN	0.43	0.79	0.56	0.877
Logistic Regression	0.24	0.85	0.37	0.943

Five-fold cross-validation of the tuned model (default 0.5 threshold per fold, hyperparameters fixed) gives $F1 = 0.784 \pm 0.029$, precision = 0.740 ± 0.066 , recall = 0.838 ± 0.028 – confirming the single-split test result is not an artifact of a favorable split, and giving a more conservative estimate than the threshold-tuned single-split number in Table 2 (the per-fold threshold was left at 0.5 deliberately, to isolate the model’s raw discriminative behavior from the threshold-tuning step).

Table 3: Comparison against prior AI4I 2020 benchmarks. Metric definitions are not identical across studies (macro-F1 averages both classes; positive-class F1 reports the failure class alone, which is substantially harder under 3.39% prevalence) – we report both where available rather than force a single number.

Study	Model	Metric	Value
Tseng [2026]	Random Forest	positive-class F1	0.722
Khemani and Qureshi [2026]	LightGBM (5-fold CV, SMOTE in-fold)	macro-F1	0.814 ± 0.012
This work	XGBoost (tuned, single split)	positive-class F1	0.90
This work	XGBoost (tuned, single split)	macro-F1	0.95
This work	XGBoost (tuned, 5-fold CV)	positive-class F1	0.784 ± 0.029

5.2 Resampling comparison and decision boundaries

Table 4 shows the validation-set bake-off; Borderline-SMOTE wins on F1 and is used throughout. Figure 2 visualizes the effect directly: a Logistic Regression classifier is fit on a 2D PCA projection of the training data, once per resampler. Without resampling, the decision boundary barely bends toward the minority class; each oversampler pulls it further into majority territory, visibly trading precision for recall – the same tradeoff Table 2 shows numerically for the full feature space and stronger models.

Table 4: Resampling technique comparison, validation set, fixed default XGBoost.

Sampler	Precision	Recall	F1	PR-AUC
Borderline-SMOTE	0.814	0.838	0.826	0.883
SMOTE	0.699	0.853	0.768	0.892
ADASYN	0.683	0.824	0.747	0.876




Figure 2: Logistic Regression decision boundary in 2D PCA space, by sampling technique. Without resampling the boundary is almost entirely majority-class; each oversampler shifts it toward the minority region.

5.3 Dimensionality reduction

The engineered features are not redundant: Figure 3 shows the PCA explained-variance elbow, and 5 of the 12 processed features are needed to reach 95% cumulative variance, so there is little to compress away by design – consistent with each engineered feature capturing a distinct mechanism (Table 1) rather than a redundant transformation of the others.




Figure 3: PCA explained-variance elbow on the processed training features. 5 of 12 components are needed for 95% cumulative variance.

Figure 4 projects the training data into 2D with PCA, t-SNE [van der Maaten and Hinton, 2008], and UMAP [McInnes et al., 2018]. PCA (linear) shows the two classes smeared together with no separating direction, consistent with linear models’ weak performance in Table 2. Both nonlinear methods instead show failure points clustering into several distinct pockets rather than one contiguous region – consistent with **Machine failure** being an aggregate of five mechanistically distinct causes rather than a single failure mode with one shared signature.

t-SNE cluster shape and spacing can be an artifact of the perplexity parameter rather than a property of the data [Wattenberg et al., 2016], so Figure 5 reruns t-SNE on the same training data at perplexity $\in \{5, 30, 50, 100\}$. The failure points consistently fall into the same right-side region of the embedding across all four settings – the pocket structure in Figure 4 is not a one-off artifact of the perplexity=30 default used there.




Figure 4: PCA vs. t-SNE vs. UMAP projections of the training data, colored by Machine failure.




Figure 5: t-SNE embeddings of the training data at increasing perplexity. Failure points (orange) occupy a consistent region of the embedding at every setting.

5.4 XAI cross-check

Table 5 reports the outcome of the protocol in Section 4.5. All five mechanisms pass: the model’s top attributed feature for each of TWF, HDF, PWF, and OSF matches that mechanism’s known physical driver, and for RNF no single feature accounts for more than 17% of total subgroup attribution (below the 0.35 threshold), consistent with the fact that random failure has no learnable cause. Figure 6 shows the full per-feature, per-mechanism attribution heatmap underlying this table, and Figure 7 shows the global SHAP summary for context.

Table 5: XAI cross-check: does the model’s top SHAP-attributed feature, restricted to rows where mechanism m fired, match m ’s known physical cause?

Mechanism	Top attributed feature	Known cause	Result
TWF	Tool wear [min]	Tool wear	✓
HDF	Rotational speed [rpm]	Temp. gap / rotational speed	✓
PWF	Power [W]	Torque $\times$ rpm	✓
OSF	Overstrain_ratio	Tool wear $\times$ torque	✓
RNF	(none dominant, top share 0.17)	none – pure chance	✓




Figure 6: Mean |SHAP value| per feature, restricted to the rows where each mechanism fired.




Figure 7: Global SHAP summary for the tuned model over the full dataset.

6 Discussion

6.1 What a failed audit would have shown

Had the model instead attributed, say, Tool wear as the dominant feature for HDF rows, that would indicate the classifier had latched onto a correlate of heat-dissipation failure (perhaps because HDF and TWF co-occur more than chance in this particular data split) rather than the actual temperature/speed mechanism – a case where the model could still score well on held-out data while being wrong about *why*, a distinction that plain accuracy or even SHAP summary plots alone would not surface. The RNF negative control is what makes the audit meaningful rather than tautological: without it, a protocol that always finds *some* top feature for every subgroup provides no evidence of discriminating validity.

6.2 Limitations

AI4I 2020 is a synthetic dataset generated from the same closed-form rules used as our ground truth (Table 1); the “ground truth” is therefore the simulator’s logic, not an independently verified physical measurement from real equipment. This is a genuine limitation for any claim that the protocol validates *real-world* causal faithfulness, and we do not make that claim here – our contribution is the protocol and its demonstration, not a validated causal finding about real machinery. This concern is not merely theoretical: [Autran et al. \[2024\]](#) show that injecting realistic irregularities (missing data, irregular time-series sampling) into AI4I 2020 measurably changes downstream model performance, underscoring that results on the clean synthetic version – ours included – may not transfer directly to messier real deployments. RNF (19 rows total dataset-wide) and TWF (46 rows) are small enough that per-mechanism subgroup SHAP estimates carry non-trivial variance; we mitigate this by computing attributions over the full dataset rather than a single split, but a larger or non-synthetic multi-mechanism dataset would give a more convincing demonstration. Finally, the audit as described requires sub-labels for at least a validation subset of failures – available here by construction, and available in many industrial settings where maintenance logs record root cause even when the model is only asked to predict binary failure, but not universal.

6.3 Generalizability

The protocol requires only two things: (1) a classifier trained on a coarse target, and (2) a held-out set of finer-grained sub-labels for at least some of the positive cases, not used in training. Both conditions are common in industrial PdM: maintenance tickets are frequently logged with a root-cause code even when the deployed model only predicts a binary failure risk. In that setting, this protocol offers a cheap, dataset-native alternative to constructing synthetic ground-truth benchmarks purely to evaluate explanation faithfulness.

7 Conclusion

We show that AI4I 2020’s five failure mechanisms, ordinarily discarded once the composite failure label is formed, can be repurposed as a faithfulness probe for post-hoc SHAP explanations. A tuned XGBoost classifier trained only on the composite label attributes each of four real mechanisms to its correct physical driver and shows no dominant feature for the one mechanism with no real cause. Beyond the specific result, we position this as a general, low-cost protocol applicable wherever a coarse PdM target aggregates several distinct, partially-labeled failure modes – the common case

in practice – as a complement to, not a replacement for, existing faithfulness-evaluation methods built on purpose-designed synthetic data.

Code and data availability. Code, the trained model artifacts, and the full notebook reproducing every table and figure in this paper are available at <https://github.com/kowshik24/glowing-broccoli>. AI4I 2020 is public [Matzka, 2020] and included in the repository.

References

- J.-V. Autran, V. Kuhn, J.-P. Diguët, M. Dubois, and C. Buche. AI4I-PMDI: Predictive maintenance datasets with complex industrial settings’ irregularities. *Procedia Computer Science*, 246:1201–1209, 2024. doi: 10.1016/j.procs.2024.09.546.
- Nitesh V. Chawla, Kevin W. Bowyer, Lawrence O. Hall, and W. Philip Kegelmeyer. Smote: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16:321–357, 2002.
- Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016.
- Shaunak Dhande, Chutian Ma, Giacinto Paolo Saggese, Paul Smith, and Krishna Taduri. A benchmark of causal vs. correlation ai for predictive maintenance. *arXiv preprint arXiv:2512.01149*, 2026.
- Hui Han, Wen-Yuan Wang, and Bing-Huan Mao. Borderline-smote: A new over-sampling method in imbalanced data sets learning. In *International Conference on Intelligent Computing*, pages 878–887. Springer, 2005.
- Haibo He, Yang Bai, Edwardo A. Garcia, and Shutao Li. Adasyn: Adaptive synthetic sampling approach for imbalanced learning. In *2008 IEEE International Joint Conference on Neural Networks*, pages 1322–1328. IEEE, 2008.
- Kushal Khemani and Anjum Nazir Qureshi. Ai-driven predictive maintenance with environmental context integration for connected vehicles: Simulation, benchmarking, and field validation. *arXiv preprint arXiv:2603.13343*, 2026.
- Yang Liu, Sujay Khandagale, Colin White, and Willie Neiswanger. Synthetic benchmarks for scientific research in explainable machine learning. *arXiv preprint arXiv:2106.12543*, 2021.
- Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30, 2017.
- Stephan Matzka. Explainable artificial intelligence for predictive maintenance applications. In *2020 Third International Conference on Artificial Intelligence for Industries (AI4I)*, pages 69–74. IEEE, 2020. doi: 10.1109/AI4I49448.2020.00023.
- Leland McInnes, John Healy, and James Melville. UMAP: Uniform manifold approximation and projection for dimension reduction. *arXiv preprint arXiv:1802.03426*, 2018.

- Chun-Jan Tseng. A physics-informed benchmarking framework for machine learning and tree-based ensembles in iiot-enabled predictive maintenance. *Sensors*, 26(16):5026, 2026. doi: 10.3390/s26165026.
- Laurens van der Maaten and Geoffrey Hinton. Visualizing data using t-sne. *Journal of Machine Learning Research*, 9(86):2579–2605, 2008.
- Simon Vollert and Andreas Theissler. Explainable predictive maintenance: A survey of current methods, challenges and opportunities. *arXiv preprint arXiv:2401.07871*, 2024.
- Martin Wattenberg, Fernanda Viégas, and Ian Johnson. How to use t-SNE effectively, 2016. URL <https://distill.pub/2016/misread-tsne/>. Distill.
- Wen Zhao, Jiawen Ding, Xueting Huang, and Yibo Zhang. Research on milling machine predictive maintenance based on machine learning and shap analysis in intelligent manufacturing environment. *arXiv preprint arXiv:2512.01205*, 2025.